

Automatic Synthesis of Random Generators for Numerically Constrained Algebraic Recursive Types^{*}

Ghiles Ziat¹, Vincent Botbol², Matthieu Dien³, Arnaud Gotlieb⁴, Martin Pépin^{1[0000-0003-1892-3017]}, and Catherine Dubois⁵

¹ Université Paris Cité, IRIF, 75205 Paris Cedex 13, France, {ghiles.ziat, martin.pépin}@irif.fr,

² Nomadic Labs, vincent.botbol@nomadic-labs.com

³ Normandie Université, UNICAEN, ENSICAEN, CNRS, GREYC, 14000 Caen, France, matthieu.dien@unicaen.fr

⁴ SIMULA RESEARCH LAB, arnaud@simula.no

⁵ Ecole Nationale Supérieure d'Informatique pour l'Industrie et l'Entreprise, Samovar, catherine.dubois@ensiie.fr

Abstract. In program verification, constraint-based random testing is a powerful technique which aims at generating random test cases that satisfy functional properties of a program. However, on recursive constrained data-structures (e.g., sorted lists, binary search trees, quadrees), and, more generally, when the structures are highly constrained, generating uniformly distributed inputs is difficult. In this paper, we present Testify: a framework in which users can define algebraic data-types decorated with high-level constraints. These constraints are interpreted as membership predicates that restrict the set of inhabitants of the type. From these definitions, Testify automatically synthesises a partial specification of the program so that no function produces a value that violates the constraints (e.g. a binary search tree where nodes are improperly inserted). Our framework augments the original program with tests that check such properties. To achieve that, we automatically produce uniform random samplers that generate values which satisfy the constraints, and verifies the validity of the outputs of the tested functions. By generating the shape of a recursive data-structure using *Boltzmann sampling* and generating evenly distributed finite domain variable values using constraint solving, our framework guarantees size-constrained uniform sampling of test cases. We provide use-cases of our framework on several key data structures that are of practical relevance for developers. Experiments show encouraging results.

^{*} this research was partially supported by the ANR PPS project ANR-19-CE48-0014 and the “DYNNET” project, co-funded by the Normandy County Council and the European Union (ERDF-ESF 2014-2020).

1 Introduction

Software Testing is one of the most widespread program verification techniques, and is also one of the most practical to implement. One interesting instance of it is Property-based Testing (PBT), where programs are tested by generating random inputs and results of the output are compared against software specifications. This technique has been extensively studied, for testing correctness [20,31], exhaustiveness [34], complexity [11] etc. However, this technique requires the developer to manually write the tests, that is the properties to be checked and the random generators. The latter can be particularly complicated to design, especially in the case of complex and constrained algebraic data structures.

In this field, *constraint-based random testing* [23] (commonly used in PBT [26,13]) is a powerful technique which aims at generating random test cases that satisfy functional properties of a program under test. By specifying a property that a program has to satisfy and by using uniformly-distributed inputs generators, it is possible to uncover subtle robustness faults that may be not be discovered otherwise. For instance, [1] explored the usage of PBT for testing a steam boiler, [28] explored its usage for wireless sensor network applications. It is worth noticing that generating inputs according to a uniform probability distribution is crucial to ensure that all the distinct program behaviours have the same chance to be triggered, even those which are the most constrained. The technique has been successfully applied in the field of unit testing for imperative programs [22] as well as various programming languages including Haskell [15], Prolog [3] and proof-assistant methodologies and tools such as Coq [32] or Isabelle [10]. Sampling constraint systems solutions according to a uniform distribution is a well-known difficult problem. Initially studied in the context of hardware testing [30], the problem has been studied in [21] and more extensively in [2]. Other random generation schemes are either not uniform, or very slow *e.g.* rejection sampling is generally uniform by construction, but fits very poorly with generation under constraints.

Recently, in [38] the authors introduced an automated framework capable of providing tests for functions that manipulate constrained values without requiring manual input from the programmer. The framework introduces a type language, with algebraic data-types, and constrained types *i.e.* types augmented with a membership predicate that is used to filter invalid representations. To verify that a function does not create invalid representations, the authors opted for a random testing approach. The main interest of the framework is that both the generators and the specifications are automatically extracted from the constraints specified by the user, which greatly alleviate the user's workload. Generators are uniform random value samplers used to provide input for functions, and specifications that are predicates that verify that a given value satisfies the constraint attached to its type, are used to check whether a function's output violates the constraint or not. Their tool is implemented as a pre-processor for OCaml programs, *i.e.* before compiling, programs are rewritten into augmented programs where a test suite has been added. During the pre-processing step, from each constrained type declaration τ is extracted a CSP p . Then, each p is

solved only once, that is to say that a characterisation of the set of solutions of p , called coverage, is calculated. Each coverage is then compiled into code which uniformly generates solutions which are then converted back into values of the type τ . However, to be able to solve a CSP only once per constrained type, the authors limit themselves to types involving a fixed number of numerical atoms (*e.g.* tuples), which automatically excludes recursive types. This makes it impractical as, for instance, in OCaml, real-world programs rely heavily on recursive data-types (lists, trees, sets, etc.).

This paper investigates the automatic synthesis of uniform pseudo-random generators, as in [38], but for recursive constrained types.

1.1 Contributions

- A programmable method to restrict the values a recursive type can take.
- An algorithm that uses Boltzmann generation and constraint solving to automatically derive uniform generators for recursive constrained types.
- An experimental study of the performances of our technique.

1.2 Outline

This paper is organised as follows: Sec.2 presents use cases of our methods on some examples of realistic code. Sec.3 defines our solving technique which mixes Boltzmann generation and global constraint solving. Sec.4 recalls some elementary notions about Boltzmann sampling and details some specifics about our use-case. Sec.5 presents our prototype and gives some details about its functioning, current capabilities and restrictions. We also give some details about our implementation and measure experimentally the performances of the generators we derive for recursive constrained types. Sec.6 describes some related work. Finally, Sec.7 summarises our work and discusses possible future improvements.

2 A Declarative Programming Approach

We propose a testing framework that allows programmers to specify constraints on recursive data structures. From these constraints, the framework extracts a Constraint Satisfaction Problem (CSP) which is solved in such a way that uniform random instances (*i.e.*, test cases) are generated. These instances are then used for testing functions in order to find defects.

2.1 Preliminaries

A pseudo-random generator g for an algebraic data-type τ is a function g of type $\mathcal{S} \rightarrow \tau$. Here, \mathcal{S} is the random state used by the pseudo random number generator. A constrained type is a pair $\langle \tau, p \rangle$, with τ an algebraic data-type and $p : \tau \rightarrow \text{bool}$ a predicate over values of type τ . The set of its inhabitants is

defined as $\{t \in \tau \mid p(t) = \text{true}\}$. A pseudo-random generator g for a constrained type $\langle \tau, p \rangle$ is a function $g : \mathcal{S} \rightarrow \tau$ s.t. $\forall s \in \mathcal{S}, p(g(s)) = \text{true}$.

Here, we face two main challenges for automating random testing of recursive data-types. First, we have to equip the developer with convenient means for specifying constraints attached to a given data-type. For example, we want to express that a list of integers is sorted or that a tree is a binary search tree (i.e., the left child node value is always smaller than the right one). Second, building an uniform random value generator for constrained recursive data-types is highly challenging. Recursive types can dynamically grow to an arbitrarily large size and, deriving generators for such types requires the resolution of a complex constraint system. In particular, we have to manage CSPs with an a-priori unknown number of variables and constraints. The grammar of Ocaml types and constraints annotations are given in Figure 7. In the following, we give two illustrative examples.

2.2 Example 1: Inserting an element into a set of integers

Let start with `list`, a recursive data-type associated to lists of integers, for which a possible type declaration is given in Fig.1. Using `list` to specify a

```
1 type list = Empty | Cons of int * list
```

Fig. 1: OCaml type declaration of lists of integers

Set data structure can easily be done using Testify, by using the annotation `[@@satisfying _]` and the `[alldiff]` constraint as illustrated by Fig. 2.

```
1 type uniquelist =  
2   | Empty  
3   | Cons of int * uniquelist [@@satisfying alldiff]
```

Fig. 2: OCaml type declaration of sets of integers using lists

We can automatically test the functions that manipulate instances of the `uniquelist` type by checking if they break the properties attached to it. For that, we have to define a generator and a specification for the corresponding type. To randomly generate instances, we first draw at random an instance of size n using Boltzmann generation (see Sec.4), then we build a CSP $(\mathcal{X}, \mathcal{D}, \mathcal{C})$ containing n finite domain variables and solve it using Path-oriented Random Testing (PRT) (see Sec.3) and eventually we build a random generator g able to produce uniformly distributed sets of size n .

A key aspect of Testify is based on the usage of *global constraints*, which are arithmetic-logic constraints holding over a non-fixed number of variables. In the example of Fig.2, we translate the declaration `[alldiff]` into a `all_different` global constraint implementation and used it to generate uniformly distributed solutions that can be used to populate test cases. For other recursive data-types, we use combination of multiple global constraints and arithmetic

constraints. Possible recursive data-types that can be implemented and tested in our framework include functions that generate and manipulate (un-)ordered lists and sets, trees, binary search trees, quadtrees, etc.

Fig.3 shows an example of a function which implements the insertion of an element within a set of integers and the code that is automatically generated for the testing of this function⁶.

```

1 let rec add (x:int) (l:uniquelist) : uniquelist =
2   match l with
3   | Empty -> Cons(x,Empty)
4   | Cons(h,tl) -> if x <> h then Cons(h,(add x tl) else l)
5
6 (* generated code*)
7 let add_test () =
8   let size = Random.int () in let rand_x = Random.int () in
9   let rand_l = unique_list size in
10  assert (alldiff_checker (add rand_x rand_l))

```

Fig. 3: Insertion of an element into a set, and the generated corresponding test

Here, testing the function means verifying that every output produced is indeed sorted (`assert (alldiff_checker (add rand_x rand_l))`). Note that we have used the return type annotation to automatically derive a (partial) specification for the function, but the generator we automatically synthesise can also be used to test any hand-written specification.

2.3 Example 2: Binary Search trees

Binary Search Trees (BST) are binary trees that additionally satisfy the following constraint: the key in each node is greater than or equal to any key stored in the left sub-tree, and less than or equal to any key stored in the right sub-tree. Stated differently, the keys in the tree must be in increasing order in a depth-first search traversal, in infix order. From this observation, we propose, using our framework, a possible OCaml declaration for BSTs illustrated in Fig.4.

```

1 type bst =
2   | Node of bst * (int[@collect]) * bst
3   | Leaf [@@satisfying fun x -> increasing x]

```

Fig. 4: Testify type annotation for binary search trees

Rather than defining global constraint for all user-declared data-types, we break the problem in two parts. On the one hand, we define or reuse known global constraints for lists, and on the other hand we define a way to browse data

⁶ The predicate `alldiff_checker` checks that the result list does not contain duplicates. It should not be mixed with the version of `alldiff` used in the type declaration which is used to generate randomly distributed solutions of that constraint

171 structures, in a certain order, by collecting the components that are subject to
 172 a global constraint. This is done using the `(int[@collect])` annotation.

173 Also, the order in which the structure is explored is crucial as it determines
 174 the order in which the variables will be passed to the global constraint. By
 175 default, a depth first order is assumed. For constructors with several arguments
 176 (e.g. `Node`), and for tuples, the order in which the traversal is made is mapped
 177 on the declaration order of the tuple component, that is in traversal order. Fig.5
 178 shows the code generated that traverses the tree.

```

1 let rec collect = function
2 | Node (a, b, c) ->
3   List.flatten [collect a; Collect.int b; collect c]
4 | Leaf -> []

```

Fig. 5: Generated collector for binary trees

179 Here, the primitive `Collect.int` is a primitive of our framework that takes
 180 an integer and builds the singleton list with this element. This way, we first
 181 visit the left sub-tree, the root and the right sub-tree. Using pre-order or post-
 182 order would give different results. This means that the constructor `Node` must
 183 be declared in the above order and, for example, the following would be in-
 184 valid: `Node of (int[@collect]) * binary_tree * binary_tree`. However, this re-
 185 striction can easily be lifted by providing an annotation which would allow the
 186 programmer to explicitly specify the traversal order. Similarly, a global anno-
 187 tation `[@bfs]` (resp. `[@dfs]`) could be used to specify that the structure must
 188 be traversed using a breadth first search (resp. depth first search). This will be
 189 studied in future work.

190 3 Constrained Type Solving

191 A *Constraint Satisfaction Problem* (CSP) is a triple $(\mathcal{X}, \mathcal{D}, \mathcal{C})$ where \mathcal{X} is a set
 192 of variables, \mathcal{D} is a function associating a finite domain (considered here as a
 193 subset of \mathcal{Z} without any loss of generality) to every variable and \mathcal{C} is a set of
 194 constraints, each of them being $\langle var(c), rel(c) \rangle$, where $var(c)$ is a tuple of
 195 variables $(X_{i_1}, \dots, X_{i_r})$ called the scope of c , and $rel(c)$ is a relation between
 196 these variables, i.e., $rel(c) \subseteq \prod_{k=1}^r D(X_{i_k})$. For each constraint c , the tuples of
 197 $rel(c)$ indicate the allowed combinations of value assignments for the variables
 198 in $var(c)$. Given a CSP $(\mathcal{X}, \mathcal{D}, \mathcal{C})$, an *assignment* is a vector (d_1, \dots, d_n) , which
 199 associates to each variable $X_i \in \mathcal{X}$ a corresponding domain value $d_i \in D(X_i)$.
 200 An assignment satisfies a constraint c if the projection of \mathcal{X} onto $var(c)$ is a
 201 member of $rel(c)$. The set of all satisfying assignments is called the solution set,
 202 noted $sol(\mathcal{C})$. A constraint c is said to be *satisfiable* if it contains at least one
 203 satisfying assignment, it is inconsistent otherwise. A CSP $(\mathcal{X}, \mathcal{D}, \mathcal{C})$ is satisfiable
 204 if it contains at least one assignment which satisfies all its constraints (i.e.,
 205 $sol(\mathcal{C}) \neq \emptyset$). A global constraint is an extension of CSPs with relations concerning
 206 a non-fixed number of variables. For instance, the $sort(Xs, Ys)$ global constraint

207 [29] takes as inputs two lists of n finite domain variables Xs, Ys (where n is
 208 unknown) and states that for each satisfying assignment $(d_1, \dots, d_n, d'_1, \dots, d'_n)$ of
 209 the constraint, we have $\forall j, \exists i$ s.t. $d'_j = \sigma(d_i)$ and $d'_1 \leq \dots \leq d'_n$, where σ is a
 210 permutation of $[1..n]$. Filtering a global constraint $c(X_1, \dots, X_n)$ with the *bound-*
 211 *consistency* local filtering property means to find D' such that for all i , the
 212 extrema values of $D'(X_i)$ are parts of satisfying assignments of c .

213 3.1 Path-Oriented Random Testing

214 *Path-oriented Random Testing* (PRT) is basically a divide-and-conquer algo-
 215 rithm, introduced in [22], which aims to generate a uniformly distributed subset
 216 of solutions of a CSP. Starting from an initial filtering step result, the general
 217 idea is to fairly divide the resulting search space into boxes of equal volumes and,
 218 after having discarded inconsistent boxes using constraint refutation, to draw at
 219 random satisfying assignments.

220 More precisely, applying constraint filtering results in domains that can be
 221 over-approximated by a larger box (i.e., an hyper-cuboid) that contains all the
 222 filtered domains. Based on an external division parameter k , PRT then fairly
 223 divides the box into k subdomains of *equal* volume. When a subdomain cannot
 224 be divided according to the division parameter k , then it is simply extended until
 225 its area can be divided. The iteration of the process leads to a fair partition of
 226 the search space into k^n subdomains where n is the number of variables of
 227 the CSP. Then constraint refutation can be used to discard (some) subdomains
 228 which are inconsistent with the rest of the CSP. As all subdomains have the
 229 same volume, it becomes possible to sample first the remaining subdomains
 230 and then, second, to randomly draw values from these subdomains. Note that,
 231 when all the subdomains are shown to be inconsistent, then the CSP is shown
 232 to be inconsistent. This contrasts with reject-based methods which will trigger
 233 assignment candidates and will reject them afterwards, without terminating in
 234 a reasonable amount of time.

Input: CSP: $(\mathcal{X}, \mathcal{D}, \mathcal{C})$, k, N : #Sol. - **Output:** t_1, \dots, t_N or \emptyset (Inconsistent)
 $\mathcal{D}' := \text{boxfilter}_{bc}(\mathcal{X}, \mathcal{D}, \mathcal{C})$; $(H_1, \dots, H_p) := \text{Fairly_Divide}(\mathcal{D}', k)$; $T := \emptyset$;
while $N > 0$ *and* $p \neq 0$ **do**
 | Pick up uniformly H at random from H_1, \dots, H_p ;
 | **if** H is inconsistent w.r.t. \mathcal{C} **then**
 | | remove H from H_1, \dots, H_p
 | **else**
 | | Pick up uniformly t at random from H and remove it;
 | | **if** \mathcal{C} is satisfied by t **then**
 | | | add t to T ; $N := N - 1$;
 | | **end**
 | **end**
end
return T ;

Algorithm 1: Path-Oriented Random Testing adapted from [22] to the uni-
 form random generation of N solutions of a CSP

235 The PRT algorithm, adapted from [22] to the case of CSP solution sam-
 236 pling, is given in Figure 1. It takes as inputs a CSP, a division parameter k ,
 237 and N a non-negative integer. Here, we make the hypothesis that, if the CSP
 238 is consistent, it contains more than N solutions. The algorithm outputs a se-
 239 quence of N uniformly distributed random assignments which satisfy the CSP.
 240 If the CSP is unsatisfiable, then PRT returns \emptyset . After an initial filtering step
 241 using bound-consistency, the algorithm partitions the resulting surrounding box
 242 in subdomains of equal volume (`Fairly_Divide` function). Then, for each lo-
 243 cally consistent subdomain H in the partition, value assignments are randomly
 244 selected and checked against the constraints of the CSP. Those which do not
 245 satisfy the constraints are simply rejected. As shown in [22], this process ensures
 246 the uniform generation of tuples in the solution space.

247 3.2 Extension with Global Constraints

248 Handling global constraints is a natural extension of PRT as it allows us to
 249 handle recursive constrained data-types. As the shape of the data structure is
 250 unknown at constraint generation time, the number of variables to be handled is
 251 also unknown in the general case. Thus, using global constraints in this context
 252 is particularly useful as it allows us to avoid the decomposition of a global con-
 253 straint into the conjunction of several simpler constraints. This results in both a
 254 stronger and faster pruning. In order to handle recursive constrained data-types,
 255 we had to provide a dedicated interface for accessing the deductions from global
 256 constraint solving. To facilitate the access to global constraints, we created an
 257 API which provides results of PRT over different global constraint combina-
 258 tions. The API provides access to predicates such as `increasing_list(+int`
 259 `LEN,+int GRAIN,-var L)` in which `L` is instantiated to a list of `LEN` uniformly
 260 distributed random integers ranked in increasing order, and the random gen-
 261 erator is initialised with `GRAIN`. Optionally, the predicate can be called with
 262 domain constraints in order to constrain the returned list of values in specific
 263 subdomains. Other similar predicates are provided as part of the API, namely
 264 `increasing_strict_list/3 (+int LEN,+int GRAIN,-var L)` which returns
 265 a list of strictly increasing integers; `decreasing_list/3` (resp. `decreasing_-`
 266 `strict_list/3`) which provides a list of integers in (resp. strict) decreasing
 267 order or else `alldiff_list/3` which returns a list of uniformly distributed ran-
 268 dom distinct integers. PRT can also be used in combination with any available
 269 global constraint and arithmetico-logic constraint. The following example, given
 270 in Fig.6 illustrates how PRT is used in this respect.

271 In this example, PRT is used with one global constraint, namely `sort(Xs, Ys)`,
 272 and some domain and arithmetic constraints to populate a constrained binary
 273 search tree (BST) of size 6. In this example, the shape of the BST is unknown
 274 and some constraints hold over the keys: the domain of the key-variables is
 275 constrained (from an externally specified source), *e.g.*, key $X_1 \in [-2.8, 2.8]$, key
 276 $X_2 \in [-3.5, 3.5]$, etc. and any key of the BST corresponds to the sum of its chil-
 277 dren (if any), *e.g.*, $Y_{father} = Y_{child_1} + Y_{child_r}$. Note that the keys have to be set
 278 in increasing order to correspond to a valid BST. Note also that we ignore in




Fig. 6: Generation of a constrained BST of size 6. Division parameter=2, length of seq.=3, 60 subdomains over 64 have been discarded after the first filtering.

MP: Il est un peu problématique cet exemple je trouve, pour deux raisons.
 (1) Il ne correspond pas au type des BST qu'on a donné plus haut où on a que des clefs et pas de valeurs, c'est pour ça que le reviewer 1 râle et ne comprends pas ce que les X_i viennent faire là.
 (2) Je comprends qu'avec PRT on peut faire des trucs sophistiqués, mais dans testify on n'a aucun moyen de spécifier ce genre de contraintes non ? Ça ne me paraît pas stratégique de parler de contraintes qu'on se sait pas exprimer.

which order will the keys be positioned in the tree. The first step of our method corresponds to the generation of a uniformly distributed random shape of the BST (Fig.6(a)) using the Boltzmann method, described in Sec.4. Then, a depth-first walk along the tree assigns variable identifiers to the nodes and collects the constraints that must hold over the constrained data structure (Fig.6(b)). The generated CSP (Fig.6(c)) can then be solved by using PRT, which generates, in this example, three uniformly distributed random solutions (Fig.6(d)(e)(f)). It is worth noticing that other uniform random solutions sampling methods such as [21,36] could have been used in this context. PRT was chosen because of its availability and simplicity. However, non-uniform random sampling such as a simple heuristic selecting at random variable and values to be enumerated first would not have been appropriate in this context as the goal was to test the robustness of user-defined functions in functional programming.

4 Boltzmann Sampling

The Boltzmann method was introduced in [17] as an algorithmic method to derive efficient sampler from *combinatorial classes*. Combinatorial classes are just sets of discrete structures with a size (a non-negative integer) and such that the number of structures having the same size is finite. For example, the binary

297 trees whose size is the number of leafs is a combinatorial class, but binary trees
 298 whose size is the length of leftmost branch is not because the number of binary
 299 trees with a leftmost branch of fixed length k is infinite. We briefly present the
 method here and refers the reader to [17] for more details.

```

<decl> ::= 'type' <type identifier> '=' <type>                                type declaration
        {'[@@satisfying' <constraints> ']}                                Testify's annotation

<type> ::= <coretype>
        | <sumtype>

<coretype> ::= 'int' | 'float' | 'char' | ...                                basic types
              | <coretype> {'*' <coretype> }                                product
              | <type identifier>

<sumtype> ::= <constructor identifier> {'[@collect]'} 'of' <coretype>
              | <sumtype> {'|' <sumtype> }

<constraints> ::= 'alldiff' | 'increasing' | 'decreasing'                    SICStus global constraints
                | <arith>                                                  arithmetic constraints like in [38, Fig. 1]
                | <constraints> {'&&' <constraints> }

```

Fig. 7: Syntax of OCaml algebraic data-types (ADT) with Testify's annotation

300

301 In the context of that paper (similarly to [12]), that method directly trans-
 302 lates into an automatic way to derive a uniform random generator of terms for
 303 the type language whose syntax is given in Fig.7. In our case, the produced
 304 generators only generate a *shape* of tree structure in a first step and the content
 305 of this shape is provided in a second step by a constraint solver which makes
 306 sure to fill the shape with values that satisfy the specified constraints. For each
 307 constrained recursive type declaration, we must therefore generate a glue func-
 308 tion between the shapes generated by the Boltzmann sampling method and the
 309 solutions returned by the solver used. This function is illustrated in the case of
 310 binary search trees in Fig.8

```

1  let rec fill_binary_tree shape solutions =
2    match shape with
3    | Label ("Node", [x1; x2; x3]) ->
4      let x1 = fill_binary_tree x1 solutions in
5      let x2 = Testify_runtime.to_int x2 solutions in
6      let x3 = fill_binary_tree x3 solutions in
7      Node (x1, x2, x3)
8    | Label ("Leaf", []) -> Leaf

```

Fig. 8: Generated function for filling the shapes for binary search trees.

311 Here, we consider types as sets of terms (the inhabitants of the type) whose
 312 size is the number of `[@collect]` values they contain. For example, using `type binary_tree`
 313 of Sec.2, the term `Node(Node(Leaf, 3, Leaf), 25, Leaf)` has size 2.

314 In the following, we denote $\Gamma\mathcal{A}_x$ a Boltzmann sampler of parameter x for the
 315 set \mathcal{A} . Such sampler produces an object $\gamma \in \mathcal{A}$ with a probability $\frac{x^{|\gamma|}}{A(x)}$ where $|\gamma|$
 316 is the size of γ and $A(x)$ is a normalizing factor called *generating series*⁷. Note
 317 that objects of the same size have the same probability to be drawn.

318 The second interest of Boltzmann samplers is that they compose well with
 319 sum, product and substitution *i.e.* the constructors of ADTs. Fig.9 shows the
 derivation of such samplers. At the end of the generating process, the object

```

1  type t = a * b (* a and b are 2 types previously defined *)
2  let gen_t x = gen_a x, gen_b x
3
4  type u = A of a | B of b
5  let gen_u x =
6    if random() < A(x)/(A(x) + B(x))
7    then gen_a x else gen_b x
8
9  type alst = Nil | Cons of a * alst (* alst(z) = 1 + z * alst(z) *)
10
11 let rec gen_alst x : alst =
12   if random() < 1/(1 - A(x))
13   then Nil else Cons(gen_a x, gen_aList A(x))

```

Fig. 9: Sampler derivation using Boltzmann

320
 321 drawn has a random size, but we see in the previous code that the choice of
 322 the parameter x influences the size. Note that we can precisely and efficiently
 323 compute x to target a size (see [6] or [33] for the details).

324 Still, the size is random. The last ingredient is to choose a parameter ϵ (which
 325 does not depend of the targeted size n) and keep only objects of size between
 326 $n - \epsilon$ and $n + \epsilon$. Thus, the size of the object is kept up to date during the
 327 generation and the generation is stopped if that size exceeds the upper bound
 328 $n + \epsilon$. At the end, the object may be smaller than $n - \epsilon$ in which case it is rejected
 329 too. However, the theory (see [17]) guarantees that the rejections cost remain
 330 relatively low, *i.e.* the cumulated size of objects sampled to obtain an object of
 331 size in the interval $[n - \epsilon, n + \epsilon]$ is in $\mathcal{O}(n)$. So the complexity of the overall
 332 process is linear in the size of the generated object.

333 An important point to mention is the case of polymorphic types. From a
 334 theoretical point of view they fit in the framework. But from a practical point of
 335 view it is hard to sample a “polymorphic value”. To deal with that limitation,
 336 the Boltzmann samplers are instantiated only for concrete types *e.g.* not for
 337 ‘a list but for int list.

⁷ The generating series $A(z)$ of a combinatorial class \mathcal{A} is defined by $A(z) = \sum_{\gamma \in \mathcal{A}} z^{|\gamma|}$

338 5 Implementation and Experiments

339 We have implemented the work presented in the previous sections in a tool
 340 available at the url <https://github.com/ghilesZ/Testify>. Our implementa-
 341 tion relies on several state-of-the-art tools. The derivation of OCaml code from
 342 annotated OCaml source files is done using the *ppx* framework, as in [37,5],
 343 which is a form of generic programming [24]. Pre-processors using *ppx* are ap-
 344 plied to source files before passing them on to the compiler. They can be seen as
 345 self-maps over abstract syntax trees. In our case, the source files are traversed to
 346 find OCaml type declarations and derive their associated generators. These gener-
 347 ators are then used to provide inputs for the functions that must be tested. We
 348 have implemented the techniques presented here for the global constraints that
 349 we have been able to identify in real data structures (BSTs, Sets, etc) namely
 350 *alldiff*, *increasing* and *decreasing* (both strict and large versions). Note that to
 351 extend our implementation, *i.e.* add a global constraint,, it is sufficient to add
 352 to the constraint solver a propagator for the said global constraint, as both the
 353 step of traversing the structure and the random generation procedure presented
 354 in section 3 being common to all types.

355 The work done by Testify is divided into two phases: the first is the pre-
 356 processing phase during which our tool collects some information on the types
 357 needed to build the generators. The second is the testing phase, where the gener-
 358 ated code is executed to produce inputs for the functions under test. Note that
 359 the pre-processing phase is performed only once while the testing phase can be
 360 triggered multiple times, each time one needs to run the tests.

361 We distinguish four kinds of types, for which we provide four different syn-
 362 thesis techniques:

- 363 – For non recursive unconstrained types (*e.g.* `int`, `float * (int * int) ...`)
 364 we determine at pre-processing time the function to be used as a generator.
 365 For that, we rely on the `qcheck` [15] library, which provides the primitives
 366 for building and composing generators.
- 367 – For non recursive constrained types (*e.g.* `int[@satisfying fun x -> x >= 0]`),
 368 we extract a single CSP which is solved once, still at pre-processing time.
 369 From this resolution is extracted a code that draws uniformly solutions of
 370 this CSP and rebuild from them a value of the corresponding type. This is
 371 the method described in [38].
- 372 – For recursive unconstrained types (*e.g.* lists, binary trees), we build samplers
 373 by using the Arbogen [19] tool. This tool implements the Boltzmann method
 374 presented in Sec.4. The tuning of the Boltzmann parameter is done at pre-
 375 processing time while the shape generation, and the conversion of this shape
 376 to a value of the targeted type is done at testing time.
- 377 – Finally, for recursive constrained types (*e.g.* sorted lists, binary search trees),
 378 the previous techniques are mixed together to produce efficient generators:
 379 first, a targeted size n is drawn, then, a shape of size n is sampled. We
 380 then browse the generated shape, collecting constrained values to build a
 381 CSP as explained in Sec.3. This CSP is then fed to the SICStus Prolog [4]

382 solver, which builds from it a generator using the *PRT* library [22]. Finally
383 we put together shapes and constrained values. All of these steps are made
384 at testing time, that is every time we have to generate a value we must solve
385 a CSP. This is arguably the bottleneck of our architecture, but experiments
386 still demonstrate the usability of our method.

387 5.1 Experiments

388 In this section, we focus on the performance of our automatically derived gener-
389 ators. We measure the generation times (in seconds) obtained with our method
390 for different constrained recursive types and by varying the size of the gen-
391 erated structure. The types we are interested in are: lists sorted in ascend-
392 ing order, association lists with unique keys, lists of pairs in ascending order
393 $((x, y) \leq (x', y) \Leftrightarrow x \leq x' \wedge y \leq y')$, binary search trees (unbalanced), functional
394 maps (key-value stores as binary search trees) and quadtrees. These types are
395 among the most frequents in the literature, and they only involve numerical
396 constraints, which Testify is able to manage.

Types	Targeted	Average	#Objects	time
increasing_list	10	8.50	2889	0.020
	100	93.95	13691	0.004
	1000	948.57	17763	0.003
	10000	9392.45	79	0.757
assoc_list	10	8.49	2534	0.023
	100	93.96	11949	0.005
	1000	947.87	13660	0.004
	10000	9406.92	76	0.786
bicollect	10	6.99	2492	0.024
	100	93.04	6418	0.009
	1000	947.73	16048	0.003
	10000	9456.85	1596	0.037
binary_tree	10	9.00	238690	0.001
	100	94.35	21214	0.001
	1000	948.00	3416	0.006
	10000	9740.00	1500	0.040
map	10	9.00	238690	0.001
	100	94.37	21208	0.001
	1000	947.08	3423	0.006
	10000	9047.00	1276	0.047
quad_tree	10	8.00	3590507	0.001
	100	93.88	228357	0.001
	1000	947.79	23548	0.002
	10000	9489.19	2191	0.027

Fig. 10: Generation time per object according to the size of the structure

397 The experience was to sample as much as possible constrained structures
398 during one minute. The results are shown in Fig.10. For each type we report the
399 size of the terms (number of `@collect` values) targeted, the average size of the

400 generated terms, the number of terms sampled and the average time to sample
 401 one term. The computer running the experiments has an Intel Core i7-6700 CPU
 402 cadenced at 3.40GHz with 8 GB of RAM.

403 As expected, at least for the tree-like types, we observe that the complexity
 404 is quite linear in the size of the sampled terms: the Boltzmann method keeps
 405 its promises and the use of a constraint solver proves to be fast enough to be
 406 used in our context. For most of these structures we manage to generate several
 407 hundred values per second, up to a certain structure size. These results prove the
 408 relevance of our method in the context of testing, as it can allow the user to fine-
 409 tune the generators to decide whether he wants to test his functions on several
 410 small structures and/or a few large ones. However, we may note that sampling
 411 of lists is much slower than sampling of trees. This is due to the fact that the
 412 Boltzmann method is not tailored for regular languages (such as list). It would
 413 probably be more efficient to use specialised algorithms for regular languages
 414 such as the one of [8].

415 6 Related Work

416 In this section we focus on related work dealing with constraint-based generation
 417 techniques. Constraint-based generation of test data has been exploited in white-
 418 box testing to produce inputs that will follow some execution paths, as well as in
 419 functional testing to generate constrained inputs. In [35], Senni applies constraint
 420 logic programming to systematically develop generators of structurally complex
 421 test data, e.g. red-black trees, in the context of Bounded-Exhaustive Testing.

422 PBT, as exemplified by Quickcheck for Haskell, has been adapted to many
 423 programming languages but also to proof assistants to test conjectures before
 424 proving them, e.g. [18,10,32,13]. In [18] restricted classes of indexed families
 425 of types are provided with surjective generators. In [13], the authors propose
 426 the FocalTest framework for testing - conditional - conjectures about functional
 427 programs and for automatically generating constrained values. In this work, CP
 428 global constraints are not used and thus FocalTest does not take benefit from
 429 the corresponding efficient filtering ad hoc procedures.

430 In the context of PBT of Erlang programs, De Angelis *et al* propose in [16]
 431 an approach to automatically derive generators of values that satisfy a given
 432 specification. Generation is performed via symbolic execution of the specifica-
 433 tion using constraint logic programming. A difference between their approach
 434 and ours is that we craft a suitable representation of a given type at static time,
 435 which is then compiled into an efficient generator. In [16], generators are built
 436 at execution time, while testing, which ultimately leads to a slower generation.
 437 The Coq plugin QuickChick helps to test Coq conjectures as soon as involved
 438 properties are executable. It allows the automatic synthesis of random genera-
 439 tors for algebraic data-types, recursive or not, and also the definition of simple
 440 inductive properties, e.g. a property specifying binary search trees whose ele-
 441 ments are between two bounds, to be turned into random generators of con-
 442 strained values [25]. The approach is narrowing-based, like in [14]. Such a binary

443 tree is built lazily while solving the constraints found in the inductive property
 444 while in Testify, the shape of the data structure is randomly chosen and then
 445 its elements are obtained by solving constraints. This tool comes with differ-
 446 ent primitives or mechanisms allowing for some flexibility in the distribution
 447 of the sampled values. For example the user can annotate the constructors of
 448 an inductive data-type with weights that are used when automatically deriving
 449 generators. Furthermore, it also produces proofs of the generators correctness.
 450 In [12], the authors adapt a Boltzmann model for random generation of OCaml
 451 algebraic data-types, possibly recursive, but not constrained. Generators are au-
 452 tomatically derived from type declarations. In [14], Claessen et al. propose an
 453 algorithm that, from a data-type definition, a constraint defined as a Boolean
 454 function and a test data size, produces random constrained values with a uni-
 455 form distribution. However the authors show that this uniformity has a high
 456 cost. They combine this perfect generator with a more efficient one based on
 457 backtracking. Limiting the class of constraints and combining it with an efficient
 458 solving process, Testify can generate constrained values with a uniform distri-
 459 bution in a reasonable time. Some work focus on the enumeration or sampling
 460 of combinatorial structures, like lambda-terms, using Boltzmann samplers [27],
 461 Prolog mechanisms [9] or both [7]. These approaches are dedicated to objects
 462 of recursive algebraic data-types with complex constraints, like typed lambda-
 463 terms, closed lambda-terms, linear lambda-terms, etc. This kind of constraints
 464 is out of reach of our tool whose objective is not only to generate constrained
 465 values but also to provide the programmer with syntactic facilities to specify
 466 them.

467 7 Conclusion

468 We have proposed in this paper a technique based on declarative programming,
 469 to derive generators of random and uniform values for constrained recursive
 470 types. We have proposed a small description language for recursive structure
 471 traversal which allows us to build a custom CSP for each term to be generated.
 472 The code we generate is efficient, and outperforms a naive generation technique
 473 based on rejection, and allows us to generate large recursive structures quickly.
 474 Starting from the constraints attached to a type, we first sample the shape of the
 475 value to generate and then build a CSP that encodes the valid representations of
 476 the terms that have this shape. Then, our tool uses the SICStus Prolog constraint
 477 solver to filter invalid representations and produce a uniform solution sampler.
 478 Our technique is integrated into the Testify framework, which embeds these
 479 generators within a fully automatic test system. The generators derived by our
 480 framework are fast enough to allow the user to run tests each time he compiles
 481 his code. This would allow him to be able to detect bugs very quickly and
 482 fix them before they become potentially harmful. However, we still have a lot
 483 of work to do to improve Testify. For example, we can extend the constraint
 484 language to be able to handle types with shape constraints (*e.g.* balanced trees).
 485 This would require adapting the Boltzmann technique to random sampling of

tree structures under constraints. Also, when dealing with a functional language, functions as values cannot be avoided: it will be necessary to have techniques for the derivation of generators for functions, and explore what kind of constrained functions (monotonic, bijective functions, etc.) appear in practice in programs. Moreover, in this paper we have only studied *tree-like* recursive data-structures. Some structures do not fit into this framework (*e.g.* graphs, doubly linked lists) and it would be interesting to see to what extent our methods adapt to these structures. Also, our current implementation tests functions by generating any random input, disregarding their body. This is naturally an important point of improvement. For example, one could imagine a static analysis of the body of the function, to conduct the input generation more precisely, and find bugs faster. Finally, our framework targets OCaml but the methods developed in this paper can be adapted to most programming languages and proof assistants.

References

1. Olfa Abdellatif-Kaddour, Pascale Thévenod-Fosse, and Hélène Waeselynck. Property-oriented testing: A strategy for exploring dangerous scenarios. In Gary B. Lamont, Hisham Haddad, George A. Papadopoulos, and Brajendra Panda, editors, *Proceedings of the 2003 ACM Symposium on Applied Computing (SAC), March 9-12, 2003, Melbourne, FL, USA*, pages 1128–1134. ACM, 2003.
2. Thomas E. Allen, Judy Goldsmith, Hayden Elizabeth Justice, Nicholas Mattei, and Kayla Raines. Uniform random generation and dominance testing for cp-nets. *J. Artif. Intell. Res.*, 59:771–813, 2017.
3. Cláudio Amaral, Mário Florido, and Vítor Santos Costa. PrologCheck - property-based testing in Prolog. In Michael Codish and Eijiro Sumii, editors, *Functional and Logic Programming - 12th International Symposium, FLOPS 2014, Kanazawa, Japan, June 4-6, 2014. Proceedings*, volume 8475 of *Lecture Notes in Computer Science*, pages 1–17. Springer, 2014.
4. Johan Andersson, Stefan Andersson, Kent Boortz, Mats Carlsson, Hans Nilsson, Thomas Sjöland, and Johan Widén. SICStus Prolog user’s manual. 1993.
5. Florent Balestrieri and Michel Mauny. Generic programming in OCaml. *Electronic Proceedings in Theoretical Computer Science*, 285:59–100, 12 2018.
6. Maciej Bendkowski, Olivier Bodini, and Sergey Dovgal. Polynomial tuning of multiparametric combinatorial samplers. In Markus E. Nebel and Stephan G. Wagner, editors, *Proceedings of the Fifteenth Workshop on Analytic Algorithmics and Combinatorics, ANALCO 2018, New Orleans, LA, USA, January 8-9, 2018*, pages 92–106. SIAM, 2018.
7. Maciej Bendkowski, Katarzyna Grygiel, and Paul Tarau. Random generation of closed simply typed λ -terms: A synergy between logic programming and Boltzmann samplers. *Theory Pract. Log. Program.*, 18(1):97–119, 2018.
8. Olivier Bernardi and Omer Giménez. A linear algorithm for the random sampling from regular languages. *Algorithmica*, 62(1–2):130–145, feb 2012.
9. Olivier Bodini and Paul Tarau. On uniquely closable and uniquely typable skeletons of lambda terms. In Fabio Fioravanti and John P. Gallagher, editors, *Logic-Based Program Synthesis and Transformation - 27th International Symposium, LOPSTR 2017, Namur, Belgium, October 10-12, 2017, Revised Selected Papers*, volume 10855 of *Lecture Notes in Computer Science*, pages 252–268. Springer, 2017.

- 533 10. Lukas Bulwahn. The new quickcheck for isabelle - random, exhaustive and symbolic
534 testing under one roof. In Chris Hawblitzel and Dale Miller, editors, *Certified*
535 *Programs and Proofs - Second International Conference, CPP 2012, Kyoto, Japan,*
536 *December 13-15, 2012. Proceedings*, volume 7679 of *Lecture Notes in Computer*
537 *Science*, pages 92–108. Springer, 2012.
- 538 11. Jacob Burnim, Sudeep Juvekar, and Koushik Sen. Wise: Automated test generation
539 for worst-case complexity. In *Proceedings of the 31st International Conference*
540 *on Software Engineering, ICSE '09*, page 463–473, New York, NY, USA, 2009.
541 Association for Computing Machinery.
- 542 12. Benjamin Canou and Alexis Darrasse. Fast and sound random generation for
543 automated testing and benchmarking in objective caml. In *Proceedings of the*
544 *2009 ACM SIGPLAN Workshop on ML, ML '09*, page 61–70, New York, NY,
545 USA, 2009. Association for Computing Machinery.
- 546 13. Matthieu Carlier, Catherine Dubois, and Arnaud Gotlieb. Focaltest: A constraint
547 programming approach for property-based testing. In José Cordeiro, Maria Virvou,
548 and Boris Shishkov, editors, *Software and Data Technologies - 5th International*
549 *Conference, ICSOFT 2010, Athens, Greece, July 22-24, 2010. Revised Selected*
550 *Papers*, volume 170 of *Communications in Computer and Information Science*,
551 pages 140–155. Springer, 2010.
- 552 14. Koen Claessen, Jonas Duregård, and Michał H Palka. Generating constrained
553 random data with uniform distribution. *Journal of functional programming*, 25,
554 2015.
- 555 15. Simon Cruanes. *QuickCheck inspired property-based testing for OCaml*. <https://github.com/c-cube/qcheck>.
556
- 557 16. Emanuele De Angelis, Fabio Fioravanti, Adrian Palacios, Alberto Pettorossi, and
558 Maurizio Proietti. *Property-Based Test Case Generators for Free*, pages 186–206.
559 09 2019.
- 560 17. Philippe Duchon, Philippe Flajolet, Guy Louchard, and Gilles Schaeffer. Boltz-
561 mann samplers for the random generation of combinatorial structures. *Combina-*
562 *torics, Probability and Computing*, 13(4-5):577–625, 2004.
- 563 18. Peter Dybjer, Qiao Haiyan, and Makoto Takeyama. Combining testing and proving
564 in dependent type theory. volume 2758, 06 2003.
- 565 19. Frederic Peschanski et al. *Arbogen, a fast uniform random generator of tree struc-*
566 *tures*. <https://github.com/fredokun/arbogen>.
- 567 20. Patrice Godefroid, Nils Klarlund, and Koushik Sen. Dart: Directed automated
568 random testing. *SIGPLAN Not.*, 40(6):213–223, June 2005.
- 569 21. Vibhav Gogate and Rina Dechter. A new algorithm for sampling CSP solutions uni-
570 formly at random. In Frédéric Benhamou, editor, *Principles and Practice of Con-*
571 *straint Programming - CP 2006, 12th International Conference, CP 2006, Nantes,*
572 *France, September 25-29, 2006, Proceedings*, volume 4204 of *Lecture Notes in Com-*
573 *puter Science*, pages 711–715. Springer, 2006.
- 574 22. Arnaud Gotlieb and Matthieu Petit. A uniform random test data generator for
575 path testing. *J. Syst. Softw.*, 83(12):2618–2626, 2010.
- 576 23. John Hughes. Quickcheck testing for fun and profit. In Michael Hanus, editor,
577 *Practical Aspects of Declarative Languages, 9th International Symposium, PADL*
578 *2007, Nice, France, January 14-15, 2007*, volume 4354 of *Lecture Notes in Com-*
579 *puter Science*, pages 1–32. Springer, 2007.
- 580 24. Ralf Lämmel and Simon Peyton Jones. Scrap your boilerplate: A practical design
581 pattern for generic programming. *SIGPLAN Not.*, 38(3):26–37, jan 2003.

- 582 25. Leonidas Lampropoulos, Zoe Paraskevopoulou, and Benjamin C. Pierce. Gen-
583 erating good generators for inductive relations. *Proc. ACM Program. Lang.*,
584 2(POPL):45:1–45:30, 2018.
- 585 26. Sophie Laplante, Richard Lassaigne, Frédéric Magniez, Sylvain Peyronnet, and
586 Michel de Rougemont. Probabilistic abstraction for model checking: An approach
587 based on property testing. *ACM Trans. Comput. Log.*, 8(4):20, 2007.
- 588 27. Pierre Lescanne. On counting untyped lambda terms. *Theor. Comput. Sci.*, 474:80–
589 97, 2013.
- 590 28. Andreas Löschner and Konstantinos Sagonas. Targeted property-based testing. In
591 *Proc. of the 26th ACM SIGSOFT International Symposium on Software Testing*
592 *and Analysis (ISSTA-17)*, pages 46–56, 07 2017.
- 593 29. Kurt Mehlhorn and Sven Thiel. Faster algorithms for bound-consistency of the
594 sortedness and the alldifferent constraint. In Rina Dechter, editor, *Principles and*
595 *Practice of Constraint Programming - CP 2000, 6th International Conference, Sin-*
596 *gapore, September 18-21, 2000, Proceedings*, volume 1894 of *Lecture Notes in Com-*
597 *puter Science*, pages 306–319. Springer, 2000.
- 598 30. Yehuda Naveh, Michal Rimón, Itai Jaeger, Yoav Katz, Michael Vinov, Eitan Mar-
599 cus, and Gil Shurek. Constraint-based random stimuli generation for hardware
600 verification. *AI Mag.*, 28(3):13–30, 2007.
- 601 31. Michal Palka, Koen Claessen, Alejandro Russo, and John Hughes. Testing an opti-
602 mising compiler by generating random lambda terms. *Proceedings - International*
603 *Conference on Software Engineering*, 01 2011.
- 604 32. Zoe Paraskevopoulou, Catalin Hritcu, Maxime Dénès, Leonidas Lampropoulos,
605 and Benjamin C. Pierce. Foundational property-based testing. In Christian Urban
606 and Xingyuan Zhang, editors, *Interactive Theorem Proving - 6th International*
607 *Conference, ITP 2015, Nanjing, China, August 24-27, 2015, Proceedings*, volume
608 9236 of *Lecture Notes in Computer Science*, pages 325–343. Springer, 2015.
- 609 33. Carine Pivoteau, Bruno Salvy, and Michele Soria. Algorithms for combinatorial
610 structures: Well-founded systems and Newton iterations. *Journal of Combinatorial*
611 *Theory, Series A*, 119(8):1711–1773, November 2012.
- 612 34. Colin Runciman, Matthew Naylor, and Fredrik Lindblad. Smallcheck and lazy
613 smallcheck automatic exhaustive testing for small values. In *Proceedings of the*
614 *First ACM SIGPLAN Symposium on Haskell*, volume 44, pages 37–48, 01 2008.
- 615 35. Valerio Senni and Fabio Fioravanti. Generation of test data structures using con-
616 straint logic programming. In Achim D. Brucker and Jacques Julliand, editors,
617 *Tests and Proofs - 6th International Conference, TAP@TOOLS 2012, Prague,*
618 *Czech Republic, May 31 - June 1, 2012. Proceedings*, volume 7305 of *Lecture Notes*
619 *in Computer Science*, pages 115–131. Springer, 2012.
- 620 36. Mathieu Vavrille, Charlotte Truchet, and Charles Prud’homme. Solution sam-
621 pling with random table constraints. In Laurent D. Michel, editor, *27th Inter-*
622 *national Conference on Principles and Practice of Constraint Programming, CP*
623 *2021, Montpellier, France (Virtual Conference), October 25-29, 2021*, volume 210
624 of *LIPICs*, pages 56:1–56:17. Schloss Dagstuhl - Leibniz-Zentrum für Informatik,
625 2021.
- 626 37. Jeremy Yallop. Practical generic programming in OCaml. pages 83–94, 01 2007.
- 627 38. Ghiles Ziat, Matthieu Dien, and Vincent Botbol. Automated Random Testing of
628 Numerical Constrained Types. In Laurent D. Michel, editor, *27th International*
629 *Conference on Principles and Practice of Constraint Programming (CP 2021)*,
630 volume 210 of *Leibniz International Proceedings in Informatics (LIPIcs)*, pages
631 59:1–59:19, Dagstuhl, Germany, 2021. Schloss Dagstuhl – Leibniz-Zentrum für In-
632 formatik.

633 A Revisions

634 A.1 To Reviewer 1

635 *The authors state that they deal with testing. However a crucial aspect of*
636 *testing seems neglected - checking whether the program produces correct*
637 *output (when run on the test data). The system, as described in l.324-*
638 *341, does not perform such checking. (BTW, how could we know which*
639 *results are to be produced by the randomly generated test data?). The*
640 *example of Fig.3 is an exception, see below.*

641 GZ-MD proposed answer: It is true that the main issue addressed in present
642 paper is that of automatically obtaining uniform random generators for con-
643 strained type. Our focus is set on these types only, and we provide tools to check
644 programs manipulating them:

- 645 – automatic synthesis of uniform samplers;
 - 646 – automatic verification that the invariants *encoded in the types* are satisfied.
- 647 We have added an explanation of this point in sec 2

648 These two contributions makes PBT simpler by alleviating the need to write
649 samplers and automatically checking the invariants that we can guess from the
650 types. Any other property a user wishes to check has to be specified.

651 AG proposed answer: As part of property-based testing where the test oracle
652 is provided by an user-specified property to check, there is a need to construct
653 uniform samplers of algebraic constrained recursive data types. Our paper is
654 focused on this issue and the underlying methodology can thus be applied in
655 various contexts. First, the case where only program crash/no-crash is checked
656 is useful in robustness testing and it requires the automatic synthesis of uniform
657 test inputs. Second, the case where one has a property to satisfy on the expected
658 output of the program is illustrated in our paper by the automatic verification
659 of type-based invariant. Both cases show that the contribution of the paper is a
660 brick that serve the purpose of property-based testing.

661 *The authors neglect one of basic principles of testing - that the test data*
662 *should include border (or extreme) cases. (This is because errors are*
663 *likely to manifest on such cases.)*

664 GZ-MD proposed answer: We agree that test data should include border
665 cases, but we consider this to be a complementary approach. Border cases such
666 as `nan`, `min_int`, `max_int`, *etc.* allow to look in places where it is known that
667 bugs like to hide. This usually allows to find more bugs, but we argue that this
668 is not the only place to search. Uniform generation, on the other hand, look
669 everywhere and test the robustness of the program. We believe that you need to
670 two approaches to achieve a good level of confidence in your program.

671 Il faut intégrer ce discours dans le corps du papier.

672 AG proposed answer: We agree that test input generation is not a monolithic
 673 process and that it should include several complementary approaches. In par-
 674 ticular, corner-based test input generation (a.k.a., boundary testing) is crucial
 675 to detect subtle faults. However, in property-based testing, several test input
 676 generators can be combined and our paper focusses on random sampling of alge-
 677 braic constrained recursive data types which is an open problem. We understand
 678 that this is not clearly stated in the paper and clarified the contribution in the
 679 revised submission.

generating inputs according to an uniform probability distribution

680 *I disagree with the claim (l.43-46) that is crucial to ensure that all the distinct program behaviours have*
the same chance to be triggered.

681 *For instance, it seems obvious that under uniform distribution of the in-*
 682 *put data various fragments of the program code are executed with various*
 683 *frequency/probability. (Sorry I cannot provide references.) This drawback*
 684 *is not present e.g. in [21], where test data for a given path in the program*
 685 *are generated.*

686 GZ-MD proposed answer: This is indeed wrong. What we meant here is that
 687 all inputs, even those that are heavily constrained have the same chance to be
 688 produced. This differs from the usual way to do property based testing where,
 689 in order to test a function with a pre-condition (a.k.a. receiving data from a
 690 constrained type), you usually generate unconstrained data and filter out the
 691 data that does not satisfy the pre-condition. This has two major drawbacks:

- 692 – if the constraint rejects a lot of inputs, either you have less tests or it has a
- 693 high impact on performance;
- 694 – this bias the distribution of the inputs towards less constrained data.

695 This second point is particularly problematic since constrained data is precisely
 696 where you need to check that the program does not mess up with your invariants.

697 AG proposed answer: It is true that uniform sampling over the program input
 698 space leads to a non-uniform sampling of the program paths. By performing
 699 random sampling over the input space, we ensure that each possible input test
 700 data has exactly the same probability to be triggered. That is the crucial property
 701 we want to guarantee. Again, it is the state-of-the-art approach used in property-
 702 based testing and a crucial requirement of robustness testing. We clarified this
 703 issue in the revised submission.

704 *It is stated (Conclusion, l.444...) that "We have proposed a small de-*
 705 *scription language for recursive structure traversal". This language is*
 706 *however not explained. It seems to appear in the examples, but the se-*
 707 *mantics is far from clear.*

708 *What is "[@satisfying]" (l.124) and "[@satisfying a]" (Fig. 2) ? What is*
 709 *"[@@satisfying ...]" and "[@collect]" (Fig.4) ?*

710 In brief, the `[@foo bar]` and `[foo bar]` syntax in OCaml are dedicated to
 711 annotate parts of the program, where `foo` is an alphanumeric identifier and `bar`
 712 is arbitrary OCaml code. They allow us to attach information to any expression,
 713 value or type declaration, etc. These annotations are by default ignored by the
 714 compiler, but they can be read by compiler extensions / pre-processors (such
 715 as the Testify tool presented in the paper), to implement extra features. AG
 716 suggested complementary answer: We have clarified the syntax/semantic of the
 717 language in the revised submission. MD : Especially, we have added the grammar
 718 of constraint annotations to the Figure 7.

719 *It looks like the paper deals with producing `_test_input_data` that satisfy*
 720 *some given constraints. The example of Fig.3 seems to be the only place*
 721 *where one considers if the `_test_results` satisfy the required constraints.*
 722 *However, this still does not mean that a test produces correct results.*

723 GZ-MD proposed answer (with revised text from AG): As explained before,
 724 we focus on test input generation in the context of robustness testing or property-
 725 based testing. Our contribution lies in the generation of algebraic recursive data
 726 types. For instance, if the type `nat` of non negative integers is defined using
 727 a `[@satisfying fun x -> x >= 0]` annotation and if a function of type `XXX`
 728 `-> nat list` is coded, then Testify will synthesise the following specification
 729 for `nat` lists:

```
730 let rec nat_list_spec xs =
731   match xs with
732   | [] -> true
733   | x :: xs -> (fun x -> x >= 0) x && nat_list_spec xs
```

734 This specification will then be used to check the output of such a function pro-
 735 ducing `nat` lists.

736 *The example of 258-278. It is absolutely unclear what is to be tested here.*
 737 *Are we going to test some procedure that checks whether its input is a*
 738 *BST of the chosen kind? No - such test data should result in both positive*
 739 *and negative outcome (and here we generate only BST's of the chosen*
 740 *kind). Are we going to test some procedure producing such BST's? No,*
 741 *we generate here such results, instead of input data for the procedure.*

742 *l.263-264 "any key of the BST corresponds to the sum of its two chil-*
 743 *dren". This condition cannot be satisfied by a finite tree (it implies that*
 744 *there are no leaves).*

745 We meant, "its children (if any)". This has been fixed.

746 *l.258... It is absolutely not clear what is the reason of additional con-*
 747 *straints ("Yfather = Ychldl + Ychldr" and the domain restrictions on*
 748 *X_i variables). Also, the X_i variables are not related to the search tree.*

MP: cf ma remarque dans la Figure 6. Je ne comprends pas non plus ce que ces contraintes viennent faire là.

The authors state that the proposed technique is "based on declarative programming" (l.442). However it seems to me that it has nothing to do with declarative programming, except for the constraints of CLP. What is used is the type definitions of the programming language dealt with, and code fragments added to programs in this language.

MP: TODO

A.2 To Reviewer 2

Sect. 3 is more problematic. Algorithm 1, which is the core of the approach, lacks a bit of explanation, e.g.: why p subdomains are returned instead of k ? What is `boxfilter_bc`? (bound-consistency is mentioned but never explained or referenced). A non-expert reader may struggle to understand how t cannot satisfy C if H is consistent w.r.t. C . Here a simple example would be useful (maybe using the problem in Fig. 6). I don't think there are problems of space, e.g., Sect. 1.1 and 1.2 can be merged into Sect. 1.

Thanks for the questions and suggestions to improve the description of Algo.1. In our setting k is a user-specified parameter and in the worst case, there are $p = 2^k$ subdomains returned by the `Fairly_divide` function. Hopefully, some subdomains can often be discarded using constraint refutation and thus $p < 2^k$. As observed by the reviewer, the `boxfilter` is parameterized by the local consistency property used in the solver and a non-expert may have hard time to understand why global consistency cannot always be achieved using a local consistency property for filtering. We have improved the description in the revised submission, according to the suggestion of reviewer 2.

Thus, using global constraints in this context is particularly useful as it allows us to constrain a non-fixed number of variables at once". But I wouldn't call using global constraint an extension, and the fact that the number of variables to be handled is unknown in general does not necessarily imply the use of global constraints.

GZ-MD proposed answer: We have clarified that the interest of global constraints is not related to expressiveness but to efficiency, *i.e.* we want to avoid the decomposition of a global constraint into a potentially large number of basic constraints.

In Sect. 6 the authors say that "CSP solving is not used in this work, the approach is narrowing based". But narrowing domains is actually a significant part of CSP solving. What is not used here is branching on variables to find one or all the solutions.

Catherine?

787

788

789

AG proposed answer: We agree with the reviewer that narrowing is a significant part of CSP solving and have corrected it in the revised submission.

790

791

Also, this work looks related to: https://link.springer.com/chapter/10.1007/978-3-030-31157-5_12 but that paper is not mentioned.

792

793

The work of De Angelis *et al* has been included in Section 6 of the revised submission.

794

795

796

797

798

799

800

801

In Sect. 7 the authors claim that "The code we generate is efficient, and outperforms a naive generation technique based on rejection, and allows us to generate large recursive structures quickly". Here comes the main weakness of the paper in my opinion: the evaluation does not compare against any other approach, including the mysterious "naive generation technique" mentioned in the conclusions. In this setting the claim that the proposed approach is efficient does not have any empirical support.

Matthieu

802

A.3 To Reviewer 3

803

804

A weak point of the paper is that the title promises more than what is actually covered by the paper.

805

806

807

A more accurate title has been given. It states that the focus of the paper is on the generation of numerically constrained **algebraic** Recursive Types (these types being described accurately by the grammar in 7).

808

809

MP: Je pense que ça ne suffit pas. Comme je le comprends, le reproche qui nous est fait c'est de dire qu'on fait des types contraints alors qu'on n'a des contraintes que numériques et pas du tout sur la forme des objets. À la place, je mettrais "Numerically Constrained Algebraic Data Types" dans le titre.

CD: J'ai aussi ajouté Random

810

811

Although the paper contains interesting examples, it lacks a precise specification of the kind of constraints and types covered by this approach.

812

813

We have enumerated in sec 5 the exhaustive list of GC we currently handle, why these, and what would be needed to extend this list.